

# Introduction to Machine Learning

Working Group “Computational Statistics” – Bernd Bischl et al.

# Splines

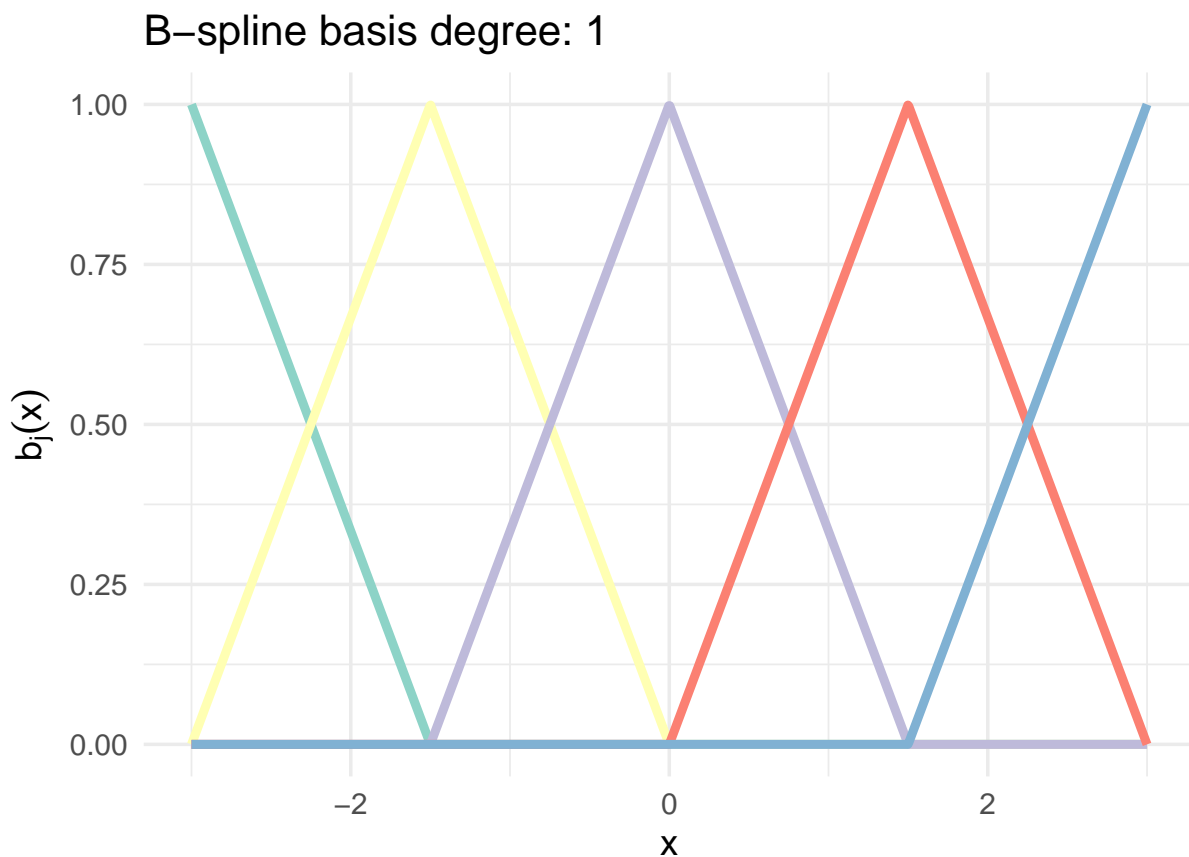
## Basic idea

A very important idea in machine learning consists in transforming features of interest or creating new features from the data. Here we will take a look at so-called splines and spline basis functions. With splines, we can model certain functional relationships arbitrarily well. To do so, we re-represent the features as evaluations of a spline basis. We can treat these new features as additional features of an *extended linear model*. Using empirical risk minimization, we can then get an estimate for the regression coefficients  $\theta_j$  associated with these spline basis features  $b_j(x), j = 1, \dots, K$ , which determine the shape of the spline function, i.e., the functional relationship  $f(x)$  between the feature  $x$  and the target  $y$ :


$$f(x) = \sum_{j=1}^K b_j(x)\theta_j,$$

## B-splines

One of the most commonly used spline bases is the so-called B-spline basis, which consists of locally defined polynomials. To use splines we have to fill our domain with knots  $t_0, \dots, t_{m+1}$ , which will serve as supporting points for the spline. For example we can use 5 equidistant knots with a piecewise polynomial degree of  $l = 1$ :



The number of B-spline basis function  $K$  is  $K = m + l + 1$ . From our machine learning perspective, this means that we are projecting the values of the feature  $x$  into a  $(3 + 1 + 1 = 5)$ -dimensional vector space with entries  $\mathbf{b}(x) = (b_1(x), b_2(x), \dots, b_5(x))$  in our little example, s.t. it holds e.g. for  $x = 0.3$ :



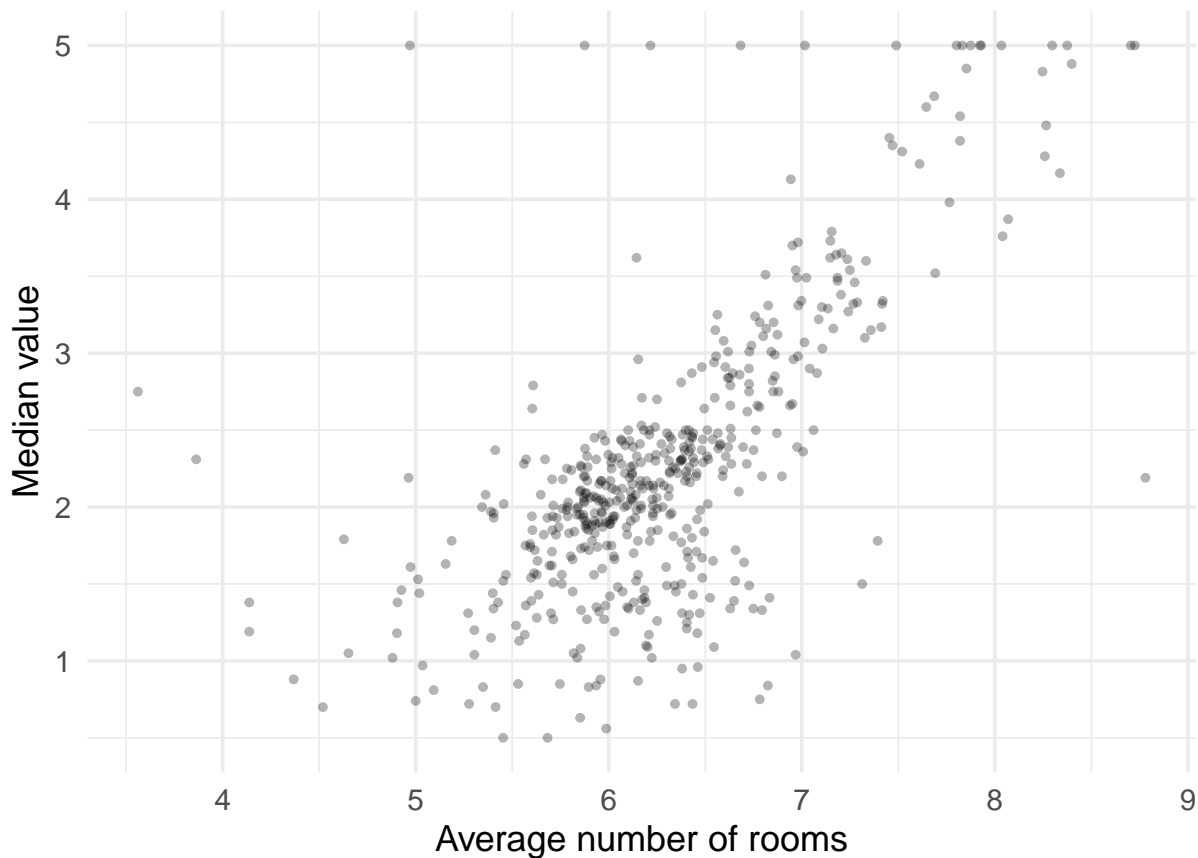
```
## [1] "b(x) = (0, 0, 0.8, 0.2, 0)"
```

## Example

Let's look at some real data, where we want to model a functional relationship:

```
library(mlbench)
library(ggplot2)
library(splines)
data(BostonHousing2)
boston_housing <- BostonHousing2
boston_housing$medv <- boston_housing$medv/10 # rescale for nicer plots

(medv_rooms_plot <- ggplot() +
  geom_point(data = boston_housing, aes(x = rm, y = medv), alpha = 0.3) +
  labs(x = "Average number of rooms",
       y = "Median value"))
```



Now let's say we want to use 5 piecewise polynomials of degree 2:

```
library(reshape2)
poly_deg <- 2
num_bfuns <- 5
num_data <- 1000 # number of points we want to use for plotting

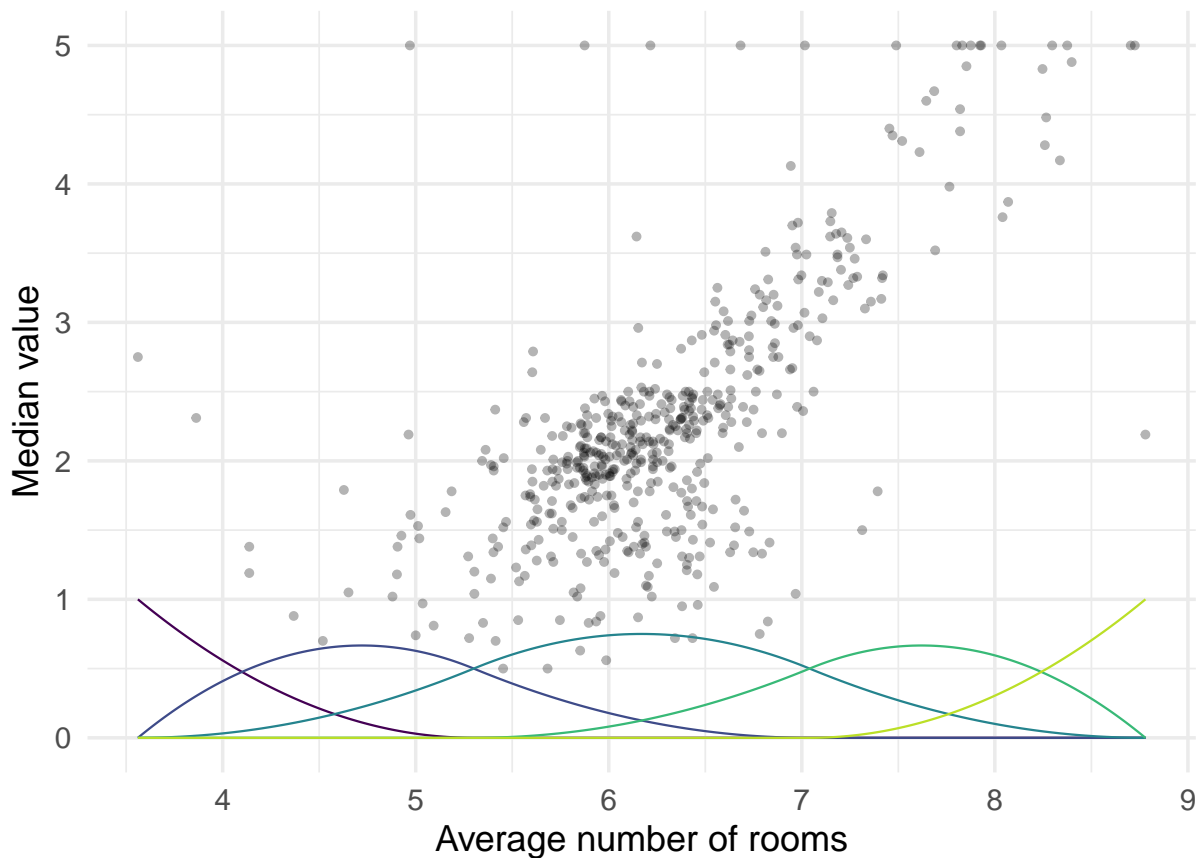
rm_min <- min(boston_housing$rm)
rm_max <- max(boston_housing$rm)

rooms <- seq(rm_min, rm_max, length.out = num_data)

bbasis_plot <- bs(rooms,
  df = num_bfuns ,
  degree = poly_deg,
  intercept = TRUE)

plot_data <- melt(data.frame(cbind(bbasis_plot, rooms)),
  id = "rooms")

medv_rooms_plot +
  geom_line(data = plot_data, aes(x = rooms, y = value, color = variable)) +
  theme(legend.position = "none")
```



Train a linear model using the 5 new features corresponding to the B-spline basis evaluations:

```
# since we want to use the same transformation, we have to specify the knots and
# piecewise polynomial degree, we have used before
bbasis_data <- data.frame(
  bs(
    x = boston_housing$rm,
    Boundary.knots = attr(bbasis_plot, "Boundary.knots"),
    knots = attr(bbasis_plot, "knots"),
    degree = poly_deg,
    intercept = TRUE
  )
)
bbasis_data$medv <- boston_housing$medv

# estimate a linear model for transformed features.
# exclude the intercept as it is contained in the B-splines basis functions.
lm_bs <- lm(medv ~ . - 1, data = bbasis_data)

# use estimated coefficients of linear model to re-scale the B-spline features:
for (i in 1:ncol(bbasis_plot)) {
  bbasis_plot[, i] = bbasis_plot[, i] *
    lm_bs$coefficients[i]
}


plot_data <-
```

```

melt(data.frame(cbind(bbasis_plot, rooms)), id = "rooms")

medv_rooms_plot +
  geom_line(data = plot_data, aes(x = rooms, y = value, color = variable)) +
  theme(legend.position = "none")

```



The estimated function relationship is simply the sum of the scaled features:


```

function_estimate <- data.frame(x = rooms)

function_estimate$y <- rowSums(bbasis_plot)

medv_rooms_plot + geom_line(data = function_estimate, aes(x = rooms, y = y)) +
  theme(legend.position = "none")

```



Wrap these steps up in one plot function:

```
# function that plots the b-spline fit
plot_bs_fit <- function(data,
                        x, y,
                        poly_deg,
                        num_bfuns,
                        x_title="x", y_title="y") {

  num_data <- 1000
  bspline <- bs(
    data[[x]],
    df = num_bfuns,
    degree = poly_deg,
    intercept = TRUE
  )
  bspline_data <- data.frame(bspline)
  bspline_data$y <- data[[y]]
  # estimate a linear model for transformed features (without intercept)
  lm_bs <- lm(y ~ . -1, data = bspline_data)

  plot_data <- data.frame(x = seq(
    min(data[[x]]),
    max(data[[x]]),
    length.out = num_data
  ))
  # scale and add up (i.e. use matrix product)
```

```

plot_data$y <- bs(
  plot_data$x,
  knots = attr(bspline, "knots"),
  degree = poly_deg,
  Boundary.knots = attr(bspline, "Boundary.knots"),
  intercept = TRUE
) %*% lm_bs$coefficients

ggplot() +
  geom_point(data = data, aes_string(x = x, y = y), alpha = 0.3) +
  geom_line(data = plot_data, aes(x = x, y = y), color = "red") +
  labs(
    title = paste(as.character(num_bfuns), "basis function(s)",
    x = x_title, #,
    y = y_title#
  )
}

```

Vary number of basis functions for cubic B-splines:

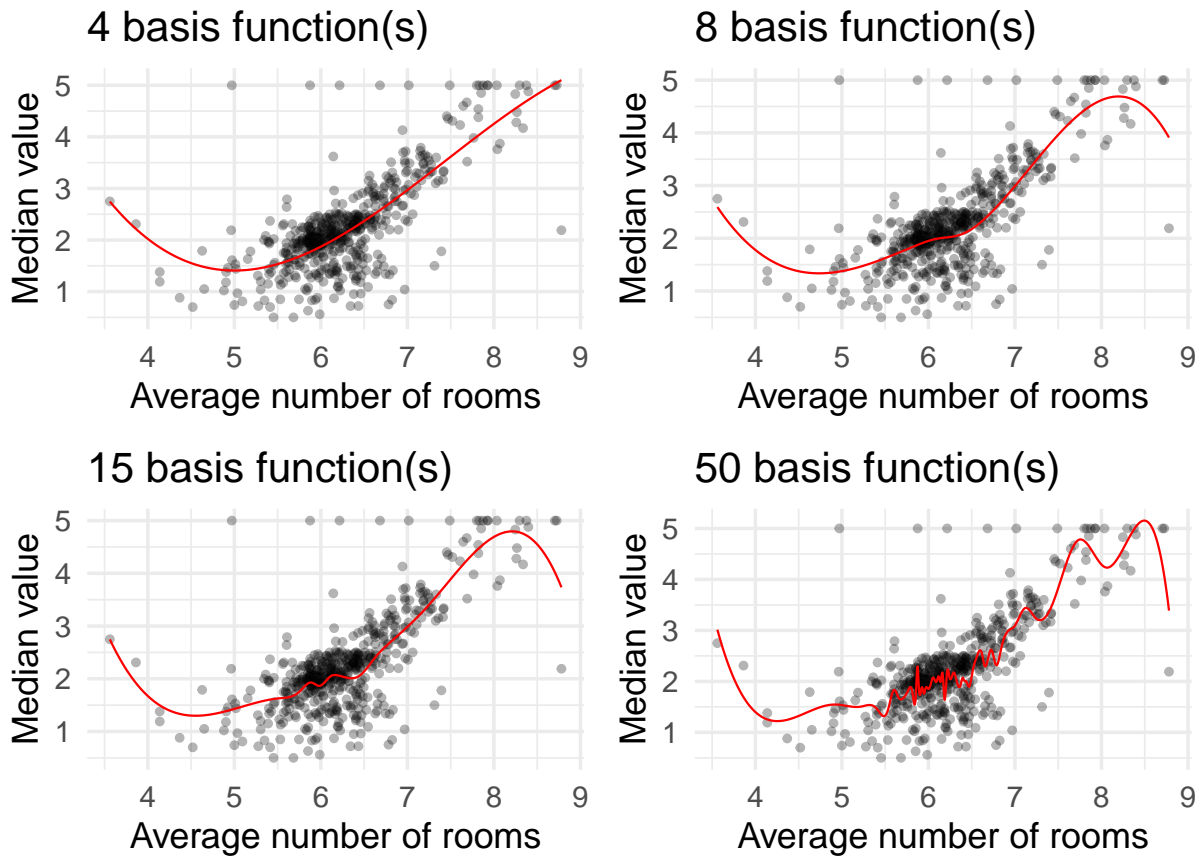
```

library(ggplot2)
library(gridExtra)

poly_deg <- 3
num_bfuns <- c(4, 8, 15, 50)
ggplot_list <- lapply(num_bfuns, function(x)
  plot_bs_fit(data = boston_housing, x = "rm", y = "medv", poly_deg = poly_deg,
    num_bfuns = x, x_title="Average number of rooms",
    y_title = "Median value"))
do.call(grid.arrange, ggplot_list)

```





- We really only fitted a simple linear model, but instead of the original feature we used transformed features.
- We observe that the fit strongly depends on the number of basis functions/knots that we choose.
- For higher number of basis functions/knots a phenomenon called overfitting, where the model fits the observed data very well, but does not generalize well on unseen data, can be seen which will be discussed in chapter 3.