Exercise 1: AdaBoost - Empirical Risk

Let $\hat{f}(\mathbf{x}) = \sum_{m=1}^{M} \hat{\beta}^{[m]} \hat{b}^{[m]}(\mathbf{x})$ be the scoring function after running AdaBoost for $M \in \mathbb{N}$ iterations. Show that the average empirical risk (on $\mathcal{D}_{\text{train}}$) of the corresponding classifier $\hat{h}(\mathbf{x}) = \text{sign}(\hat{f}(\mathbf{x}))$ is bounded as follows

$$\frac{\mathcal{R}_{\text{emp}}(\hat{h})}{n} = \frac{\sum_{i=1}^{n} \mathbb{1}_{[\hat{h}(\mathbf{x}^{(i)}) \neq y^{(i)}]}}{n} \leq \prod_{m=1}^{M} \sqrt{1 - 4(\gamma^{[m]})^2}, \tag{1}$$

where $\gamma^{[m]} = \frac{1}{2} - \text{err}^{[m]}$. For this purpose, proceed as follows:

(a) Give an interpretation of $\gamma^{[m]}$.

(b) For any $m = 1, \ldots, M$ let $W^{[m]} = \sum_{i=1}^{n} w^{[m](i)}$ be the total weigh in iteration **before** normalizing the weights. Show that $W^{[m]} = \sqrt{1 - 4(\gamma^{[m]})^2}$.

(c) Show that

$$w^{[M+1](i)} = \frac{w^{[1](i)} \exp(-y^{(i)} \hat{f}\left(\mathbf{x}^{(i)}\right))}{\prod_{m=1}^{M} W^{[m]}},$$

where $w^{[M+1](i)}$ is the **normalized** weight if we would run AdaBoost for $M+1$ iterations.

(d) Argue that $\mathbb{1}_{[\hat{h}(\mathbf{x}) \neq y]} \leq \exp(-y\hat{f}(\mathbf{x}))$ for any $(\mathbf{x}, y) \in \mathcal{X} \times \mathcal{Y}$.

(e) Combine everything to conclude (1).

Hint: Since for any x it holds that $1 + x \leq \exp(x)$, we can infer from (1) that

$$\frac{\mathcal{R}_{\text{emp}}(\hat{h})}{n} \leq \exp\left(-2 \sum_{m=1}^{M} (\gamma^{[m]})^2\right) = \exp\left(-2 \sum_{m=1}^{M} (\frac{1}{2} - \text{err}^{[m]})^2\right),$$

i.e., the average empirical risk is decreasing exponentially in the number of used iterations (provided $\text{err}^{[m]} > 1/2$).