

Solution 1: Risk Minimization and Gradient Descent

- (a) • Hypothesis space \mathcal{H} is defined as:

$$\mathcal{H} = \{f(\mathbf{x}) = \mathbf{x}^\top \boldsymbol{\beta} \mid \boldsymbol{\beta} \in \mathbb{R}^p\}$$

- We fit a linear model, ergo using the $L2$ loss makes sense (e.g., because of the link to Gaussian MLE):

$$L(y^{(i)}, f(\mathbf{x}^{(i)} | \boldsymbol{\beta})) = L(y^{(i)}, \mathbf{x}^{(i)\top} \boldsymbol{\beta}) = 0.5 \left(y^{(i)} - \mathbf{x}^{(i)\top} \boldsymbol{\beta}\right)^2$$

and the theoretical risk is

$$\mathcal{R}(f) = \mathcal{R}(\boldsymbol{\beta}) = \int L(y, f(\mathbf{x})) d\mathbb{P}_{xy} = 0.5 \int (y - f(\mathbf{x}))^2 d\mathbb{P}_{xy} = 0.5 \int (y - \mathbf{x}^\top \boldsymbol{\beta})^2 d\mathbb{P}_{xy}.$$

- (b) The Bayes regret is $\mathcal{R}_L(\hat{f}) - \mathcal{R}_L^*$ and can be decomposed into an estimation error $[\mathcal{R}_L(\hat{f}) - \inf_{f \in \mathcal{H}} \mathcal{R}_L(f)]$ and an approximation error $[\inf_{f \in \mathcal{H}} \mathcal{R}_L(f) - \mathcal{R}_L^*]$.

- (i) If $f^* \in \mathcal{H}$, $\mathcal{R}_L^* = \inf_{f \in \mathcal{H}} \mathcal{R}_L(f)$, i.e., the approximation error is 0 and for $n \rightarrow \infty$ the Bayes regret $\rightarrow 0$.
(ii) If $f^* \notin \mathcal{H}$, the Bayes regret typically consists of both parts, but as $n \rightarrow \infty$, we are left with the approximation error.

- (c) • The empirical risk is

$$\mathcal{R}_{emp}(\boldsymbol{\beta}) = 0.5 \sum_{i=1}^n \left(y^{(i)} - \mathbf{x}^{(i)\top} \boldsymbol{\beta}\right)^2 = 0.5 \|\mathbf{y} - \mathbf{X}\boldsymbol{\beta}\|^2.$$

- Optimization = minimization of the empirical risk can either be done analytically (the preferred solution in this case!) or using, e.g., gradient descent.

$$\nabla_{\boldsymbol{\beta}} \mathcal{R}_{emp}(\boldsymbol{\beta}) = 0.5 \nabla_{\boldsymbol{\beta}} (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta}) = -\mathbf{X}^\top (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$$

- (d) For convex objectives, every local minimum corresponds to a global minimum. To show convexity, calculate the second derivatives:

$$\nabla_{\boldsymbol{\beta}\boldsymbol{\beta}^\top} \mathcal{R}_{emp}(\boldsymbol{\beta}) = \mathbf{X}^\top \mathbf{X}.$$

Since $\mathbf{z}^\top \mathbf{X}^\top \mathbf{X} \mathbf{z}$ is the inner product of a vector $\tilde{\mathbf{z}} = \mathbf{X} \mathbf{z}$ with itself, i.e.

$$\mathbf{z}^\top \mathbf{X}^\top \mathbf{X} \mathbf{z} = \tilde{\mathbf{z}}^\top \tilde{\mathbf{z}} = \sum_{i=1}^n \tilde{z}_i^2$$

it is ≥ 0 and hence $\mathbf{X}^\top \mathbf{X}$ psd and therefore $\mathcal{R}_{emp}(\boldsymbol{\beta})$ convex.